How the US pipeline classifies AI statements in company filings

Run us-2025h1_V3: 7,794 US filings, January to June 2025, taxonomy v0.9

Step	Counts of this run	Where it runs
1 Collect the filings US annual-report filings (10-K, ARS, 20-F and others) filed between January and June 2025, downloaded as text from the data provider.	7,794 filings	Data provider's API → Google Cloud Storage
2 Read each filing The filing is split into its paragraphs, headings and table rows, and the report sections (Risk Factors, Business, …) are recognised.	7,794 filings read	Pipeline software, no model
3 Find the passages that may be about AI A list of 502 AI terms plus terms of the company's industry is searched in the text. Each hit and its neighbouring paragraphs form one passage.	28,378 passages with an AI term from 5,368 filings	Pipeline software, no model
Passages with context terms only (data centre, analytics, …): kept and counted, not classified: 64,643 passages. Filings without any candidate passage: 2,426.		
4 Classify each passage Gemini 2.5 Flash reads 40 passages per request together with the codebook (taxonomy v0.9: definitions and criteria for every label) and returns one label per passage: is it about AI; its frame; AI's role; risk types; MIT risk subdomains; adoption stage; business function; governance response; kind of AI; disclosure traits; substantiveness; one to three verbatim quotes as evidence.	28,378 passages sent about USD 34 of model calls	Gemini on Google Cloud (Vertex AI)
5 Check every label Each label is checked automatically: allowed values, field rules, and every quote must appear word for word in the passage. A failed label is asked again.	27,908 labelled (98.3 %) 24,765 about AI · 3,143 not about AI 470 without label (1.7 %)	Pipeline software, no model
Not about AI: the AI term was used in another sense (a product name, an abbreviation). These 3,143 passages are dropped from the AI counts.		
6 Add up per filing The passage labels are turned into statements about the filing: mentions AI, states an AI risk, reports its own use of AI, names a kind of AI, …	4,924 filings mention AI (63 %) 3,999 state an AI risk 2,696 report their own use of AI 2,870 have no AI statement	Pipeline software, no model
7 Publish the dataset Passages, labels, evidence quotes and filing metrics are written as one dataset that the analyses and the human validation read.	39,579 evidence quotes human validation: 290 passages, 7 raters	Google Cloud Storage (run us-2025h1_V3)